

Tools for OpenMP Programming

■ Correctness Tools

- ThreadSanitizer
- Intel Inspector XE (or whatever the current name is)

■ Performance Analysis

- Performance Analysis basics
- Overview on available tools

- Data Race: the typical OpenMP programming error, when:
 - two or more threads access the same memory location, and
 - at least one of these accesses is a write, and
 - the accesses are not protected by locks or critical regions, and
 - the accesses are not synchronized, e.g. by a barrier.
- Non-deterministic occurrence: e.g. the sequence of the execution of parallel loop iterations is non-deterministic
 - In many cases *private* clauses, *barriers* or *critical regions* are missing
- Data races are hard to find using a traditional debugger

ThreadSanitizer: Overview

- Correctness checking for threaded applications
- Integrated in clang and gcc compiler
- Low runtime overhead: 2x – 15x
- Used to find data races in browsers like Chrome and Firefox

ThreadSanitizer: Usage

module load clang

Module in Aachen.

<https://pruners.github.io>



Compile the program with clang compiler:

```
clang -fsanitize=thread -fopenmp -g myprog.c -o myprog
```

```
clang++ -fsanitize=thread -fopenmp -g myprog.cpp  
-o myprog
```

```
gfortran -fsanitize=thread -fopenmp -g myprog.f -c
```

```
clang -fsanitize=thread -fopenmp -lgfortran myprog.o  
-o myprog
```

- Execute:

```
OMP_NUM_THREADS=4 ./myprog
```

- Understand and correct the detected threading errors

ThreadSanitizer: Example

```
1 #include <stdio.h>
2
3 int main(int argc, char **argv) {
4     int a = 0;
5     #pragma omp parallel
6     {
7         if (a < 100) {
8             #pragma omp critical
9             a++;
10        }
11    }
12 }
```

WARNING: ThreadSanitizer: data race

Read of size 4 at 0x7ffffffdcdc by thread T2:

#0 .omp_outlined. race.c:7
(race+0x0000004a6dce)
#1 __kmp_invoke_microtask <null>
(libomp_tsan.so)

Previous write of size 4 at 0x7ffffffdcdc by
main thread:

#0 .omp_outlined. race.c:9
(race+0x0000004a6e2c)
#1 __kmp_invoke_microtask <null>
(libomp_tsan.so)

■ Detection of

- Memory Errors
- Deadlocks
- Data Races

■ Support for

- WIN32-Threads, Posix-Threads, Intel Threading Building Blocks and OpenMP

■ Features

- Binary instrumentation gives full functionality
- Independent stand-alone GUI for Windows and Linux


PI example / 1

```
double f(double x)
{
    return (4.0 / (1.0 + x*x));
}
```

```
double CalcPi (int n)
{
    const double fH = 1.0 / (double) n;
    double fSum = 0.0;
    double fX;
    int i;
```

```
#pragma omp parallel for private(fX,i) reduction(+:fSum)
    for (i = 0; i < n; i++)
    {
        fX = fH * ((double)i + 0.5);
        fSum += f(fX);
    }
    return fH * fSum;
}
```

$$\pi = \int_0^1 \frac{4}{1+x^2}$$



PI example / 2

```
double f(double x)
{
    return (4.0 / (1.0 + x*x));
}

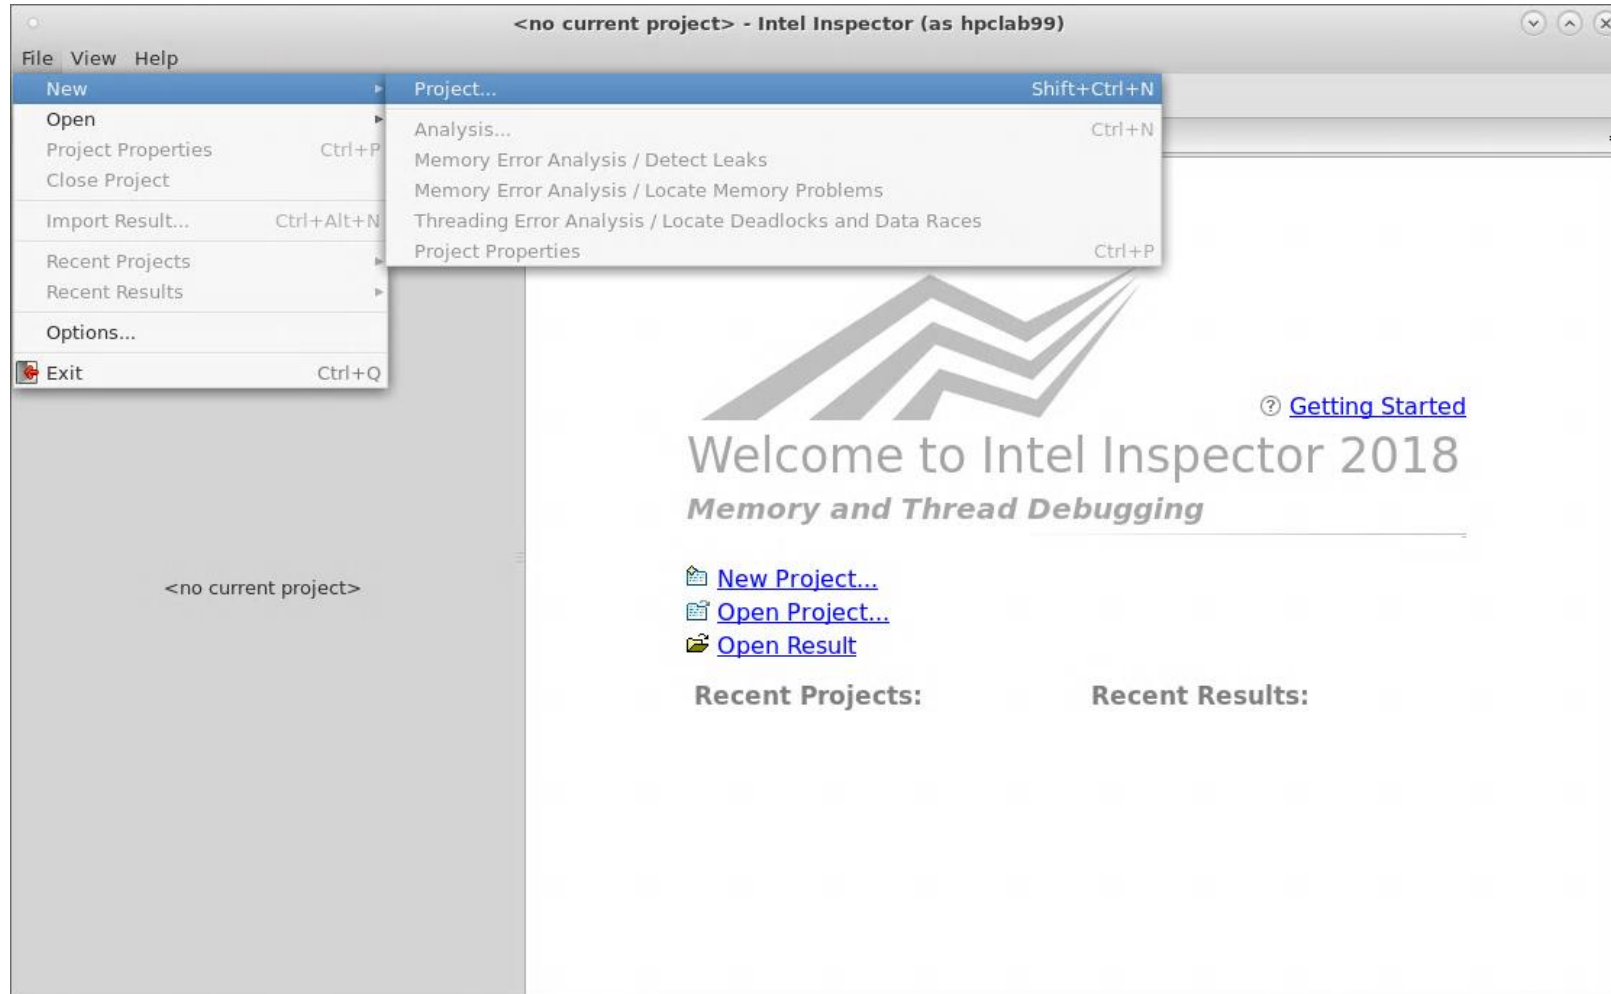
double CalcPi (int n)
{
    const double fH = 1.0 / (double) n;
    double fSum = 0.0;
    double fX;
    int i;

    #pragma omp parallel for private(fX,i) reduction(+:fSum)
    for (i = 0; i < n; i++)
    {
        fX = fH * ((double)i + 0.5);
        fSum += f(fX);
    }
    return fH * fSum;
}
```

What if we
would have
forgotten this?

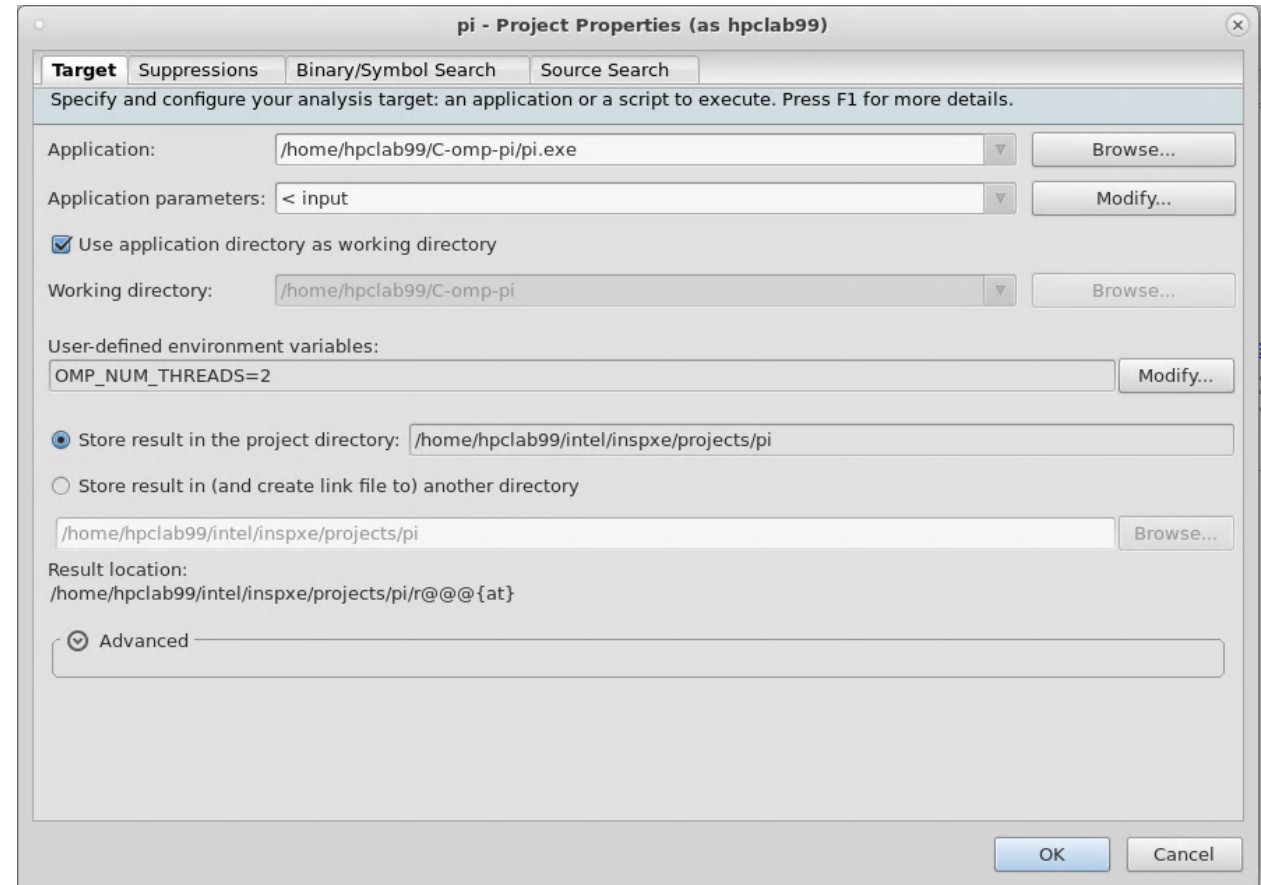
Inspector XE: create project / 1

```
$ module load Inspector ; inspxe-gui
```



Inspector XE: create project / 2

- ensure that multiple threads are used
- choose a small dataset (really!), execution time can increase 10X – 1000X

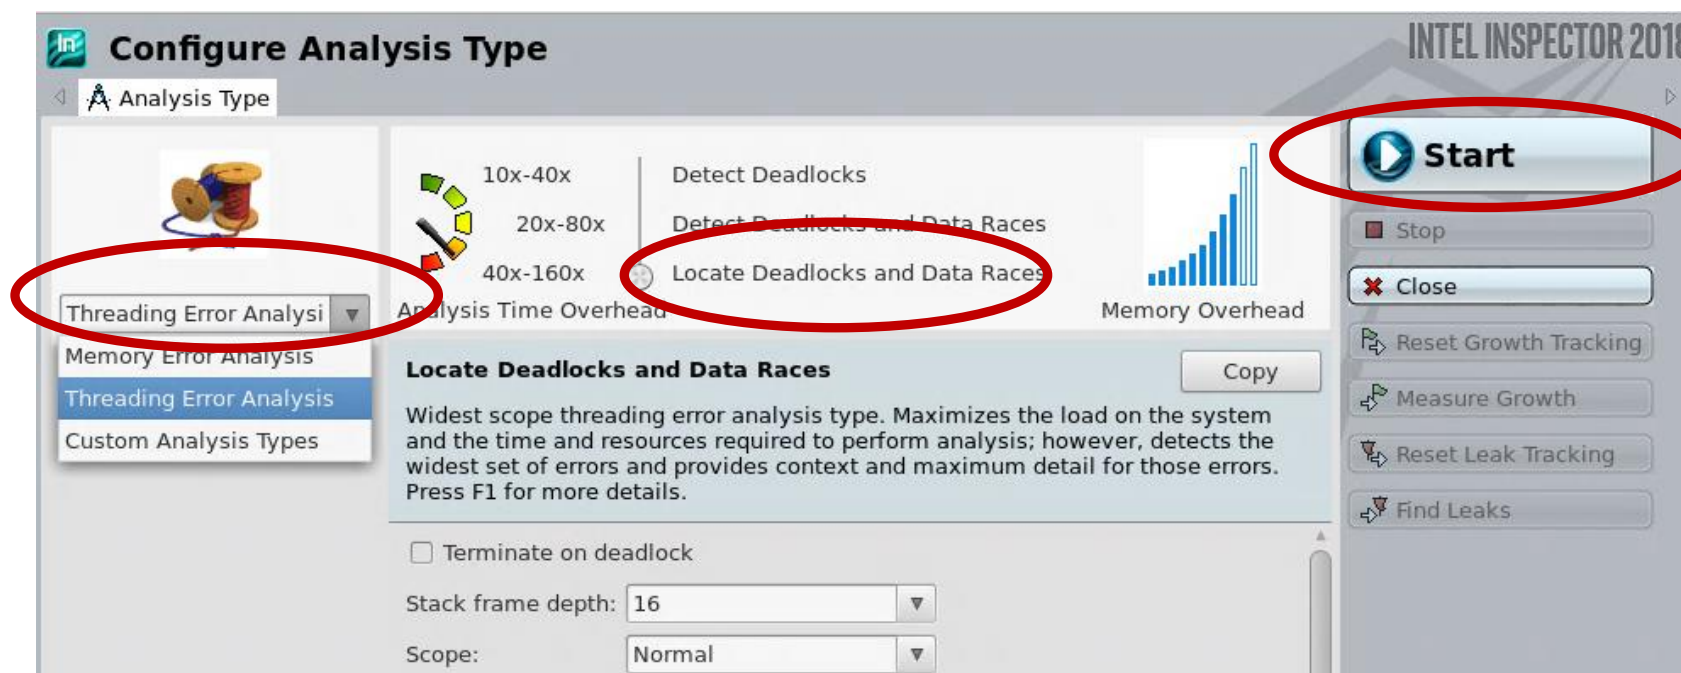


Inspector XE: configure analysis

Threading Error Analysis Modes

1. Detect Deadlocks
2. Detect Deadlocks and Data Races
3. Locate Deadlocks and Data Races

more details,
more overhead



Inspector XE: results / 1

- 1 detected problems
- 2 filters
- 3 code location
- 4 Timeline

The screenshot displays the Intel Inspector 2018 interface. The main window is titled "/home/hpclab99/intel/inspxe/projects/pi - Intel Inspector (as hpclab99)". The "Locate Deadlocks and Data Races" section is active, showing a list of detected problems. A yellow circle '1' highlights the "Problems" table, which lists three data race issues. A yellow circle '2' highlights the "Filters" panel on the right, which shows filters for Severity (Error), Type (Data race), Source (pi.c), Module (pi.exe), and State (New). Below the problems table, a yellow circle '3' highlights the "Code Locations: Data race" section, which shows the source code for the detected data race. A yellow circle '4' highlights the "Timeline" section, which shows the execution timeline of the program, including the OMP Master Thread #0 (23581) and OMP Worker Thread #1 (23717).

ID	Type	Sources	Modules	State
P1	Data race	pi.c	pi.exe	New
	Data race	pi.c:72	pi.exe	New
	Data race	pi.c:72	pi.exe	New

Description	Source	Function	Module	Variable
Read	pi.c:72	CalcPi	pi.exe	
<pre>70 { 71 fX = fH * ((double)i + 0 72 fSum += f(fX); 73 } 74 return fH * fSum;</pre>				
Write	pi.c:72	CalcPi	pi.exe	
<pre>70 { 71 fX = fH * ((double)i + 0 72 fSum += f(fX); 73 } 74 return fH * fSum;</pre>				

Timeline: OMP Master Thread #0 (23581), OMP Worker Thread #1 (23717)

Inspector XE: results / 2

- 1 Source Code producing the issue – double click opens an editor
- 2 Corresponding Call Stack

The image displays the Intel Inspector 2018 interface for a data race analysis. The main window is titled "Data race" and shows two threads involved in a race condition. The top thread is "Read - Thread OMP Master Thread #0 (23581) (pi.exe!CalcPi - pi.c:72)" and the bottom thread is "Write - Thread OMP Worker Thread #1 (23717) (pi.exe!CalcPi - pi.c:72)". Both threads are shown with their source code and a call stack. The source code for both threads is identical, showing a loop that calculates the value of pi. The call stack for both threads shows the function "pi.exe!CalcPi" and the entry point "pi.exe!_start".

Thread 1: Read - Thread OMP Master Thread #0 (23581) (pi.exe!CalcPi - pi.c:72)

Source Code (pi.c):

```
67 //#pragma omp parallel for private(i, fX) reduction(+:fSum)
68 #pragma omp parallel for private(i, fX)
69   for (i = iRank; i < n; i += iNumProcs)
70   {
71       fX = fH * ((double)i + 0.5);
72       fSum += f(fX);
73   }
74   return fH * fSum;
75 }
76
```

Call Stack:

- pi.exe!CalcPi - pi.c:72
- pi.exe!CalcPi - pi.c:68
- pi.exe!_start

Thread 2: Write - Thread OMP Worker Thread #1 (23717) (pi.exe!CalcPi - pi.c:72)

Source Code (pi.c):

```
67 //#pragma omp parallel for private(i, fX) reduction(+:fSum)
68 #pragma omp parallel for private(i, fX)
69   for (i = iRank; i < n; i += iNumProcs)
70   {
71       fX = fH * ((double)i + 0.5);
72       fSum += f(fX);
73   }
74   return fH * fSum;
75 }
76
```

Call Stack:

- pi.exe!CalcPi - pi.c:72

Inspector XE: results / 3

- 1 Source Code producing the issue – double click opens an editor
- 2 Corresponding Call Stack

The missing reduction is detected.

Data race

Target Analysis Type Collection Log Summary Sources

Read - Thread OMP Master Thread #0 (23581) (pi.exe!CalcPi - pi.c:72)

pi.c Disassembly (pi.exe!0x111f)

```
67 // #pragma omp parallel for private(i, fX) reduction(+:fSum)
68 #pragma omp parallel for private(i, fX)
69 for (i = iRank; i < n; i += iNumProcs)
70 {
71     fX = fH * ((double)i + 0.5);
72     fSum += f(fX);
73 }
74 return fH * fSum;
75 }
76
```

1

Call Stack

- pi.exe!CalcPi - pi.c:72
- pi.exe!CalcPi - pi.c:68
- pi.exe!_start

2

Write - Thread OMP Worker Thread #1 (23717) (pi.exe!CalcPi - pi.c:72)

pi.c Disassembly (pi.exe!0x1395)

```
67 // #pragma omp parallel for private(i, fX) reduction(+:fSum)
68 #pragma omp parallel for private(i, fX)
69 for (i = iRank; i < n; i += iNumProcs)
70 {
71     fX = fH * ((double)i + 0.5);
72     fSum += f(fX);
73 }
74 return fH * fSum;
75 }
76
```

1

Call Stack


- pi.exe!CalcPi - pi.c:72

2

Sampling vs. Instrumentation


Sampling

- Running program is periodically interrupted to take measurement
- *Statistical* inference of program behavior
- Works with unmodified executables



Instrumentation


- Every event of interest is captured directly
- More detailed and *exact* information
- Typically: recompile for instrumentation



Tracing vs. Profiling

Trace

- Chronologically ordered sequence of event records




Profile from instrumentation

- Aggregated information



Profile from sampling



OMPT support for sampling


- OMPT defines states like *barrier-wait*, *work-serial* or *work-parallel*

- Allows to collect OMPT state statistics in the profile
- Profile break down for different OMPT states

```
void foo() {}  
void bar() {foo();}  
void baz() {bar();}  
int main()  
{foo();bar();baz();  
 return 0;}
```

- OMPT provides frame information

- Allows to identify OpenMP runtime frames.
- Runtime frames can be eliminated from call trees



OMPT support for instrumentation

- OMPT provides event callbacks
 - Parallel begin / end
 - Implicit task begin / end
 - Barrier / taskwait
 - Task create / schedule
- Tool can instrument those callbacks
- OpenMP-only instrumentation might be sufficient for some use-cases

```
void foo() {}  
void bar() {  
    #pragma omp task  
    foo();}  
void baz() {  
    #pragma omp task  
    bar();}  
int main() {  
    #pragma omp parallel sections  
    {foo();bar();baz();}  
    return 0;}  
}
```

VI-HPS Tools / 1

- Virtual institute – high productivity supercomputing
- Tool development
- Training:
 - VI-HPS/PRACE tuning workshop series
 - SC/ISC tutorials
- Many performance tools available under vi-hps.org
 - → tools → VI-HPS Tools Guide
 - Tools-Guide: flyer with a 2 page summary for each tool

Data collection

- Score-P : instrumentation based profiling / tracing
- Extrae : instrumentation based profiling / tracing

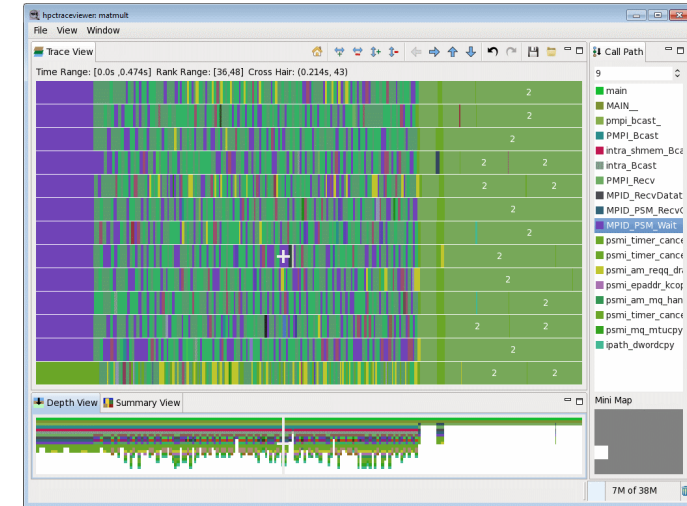
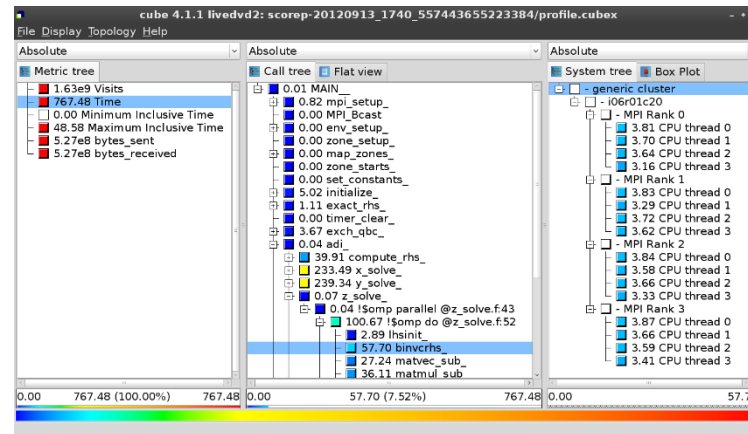

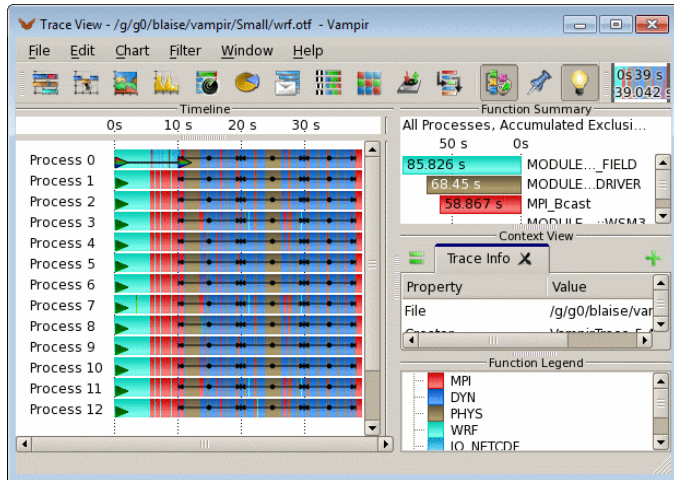
Data processing

- Scalasca : trace-based analysis

Data presentation

- ARM Map, ARM performance report
- CUBE : display for profile information
- Vampir : display for trace data (commercial/test)
- Paraver : display for extrae data
- Tau : visualization

Performance tools GUI



Correctness:

- Data Races are very hard to find, since they do not show up every program run.
- Intel Inspector XE or ThreadSanitizer help a lot in finding these errors.
- Use really small datasets, since the runtime increases significantly.


Performance:

- Start with simple performance measurements like hotspots analyses and then focus on these hot spots.
- In OpenMP applications analyze the waiting time of threads. Is the waiting time balanced?
- Hardware counters might help for a better understanding of an application, but they might be hard to interpret.


OpenMP Parallel Loops

- Existing loop constructs are tightly bound to execution model:


```
#pragma omp parallel for  
for (i=0; i<N;++i) {...}
```



```
#pragma omp simd  
for (i=0; i<N;++i) {...}
```



```
#pragma omp taskloop  
for (i=0; i<N;++i) {...}
```



- The `loop` construct is meant to tell OpenMP about truly parallel semantics of a loop.

OpenMP Fully Parallel Loops

```
int main(int argc, const char* argv[]) {  
    float *x = (float*) malloc(n * sizeof(float));  
    float *y = (float*) malloc(n * sizeof(float));  
    // Define scalars n, a, b & initialize x, y  
  
    #pragma omp parallel  
    #pragma omp loop  
        for (int i = 0; i < n; ++i){  
            y[i] = a*x[i] + y[i];  
        }  
    }  
}
```

Loop Constructs, Syntax

■ Syntax (C/C++)

```
#pragma omp loop [clause[[, clause],...]  
for-loops
```

■ Syntax (Fortran)

```
!$omp loop [clause[[, clause],...]  
do-loops  
[!$omp end loop]
```

Loop Constructs, Clauses

- `bind(binding)`

- Binding region the loop construct should bind to
- One of: `teams`, `parallel`, `thread`

- `order(concurrent)`

- Tell the OpenMP compiler that the loop can be executed in any order.
- Default!

- `collapse(n)`

- `private(list)`

- `lastprivate(list)`

- `reduction(reduction-id: list)`

Extensions to Existing Constructs

- Existing loop constructs have been extended to also have truly parallel semantics.

- C/C++ Worksharing:

```
#pragma omp [for|simd] order(concurrent) \  
                                     [clause[[,] clause],...]  
  
for-loops
```

- Fortran Worksharing:

```
!$omp [do|simd] order(concurrent) &  
                                     [clause[[,] clause],...]  
  
do-loops  
[!$omp end [do|simd]]
```

DOACROSS Loops

- “DOACROSS” loops are loops with special loop schedules
 - Restricted form of loop-carried dependencies
 - Require fine-grained synchronization protocol for parallelism


- Loop-carried dependency:
 - Loop iterations depend on each other
 - Source of dependency must be scheduled before sink of the dependency

- DOACROSS loop:
 - Data dependency is an invariant for the execution of the whole loop nest

Parallelizable Loops

- A parallel loop cannot not have any loop-carried dependencies (simplified just a little bit!)


```
for (int i = 1; i < N; ++i) {  
    for (int j = 1; j < M; ++j) {  
        b[i][j] = f(b[i][j],  
                    b[i][j], a[i][j]);  
    }  
}
```



Non-parallelizable Loops

- If there is a loop-carried dependency, a loop cannot be parallelized anymore (“easily” that is)


```
for (int i = 1; i < N; ++i) {  
    for (int j = 1; j < M; ++j) {  
        b[i][j] = f(b[i-1][j],  
                   b[i][j-1], a[i][j]);  
    }  
}
```



Wavefront-Parallel Loops

- If the data dependency is invariant, then skewing the loop helps remove the data dependency

```
for (int i = 1; i < N; ++i) {  
    for (int j = i+1; j < i+N; ++j) {  
        b[i][j-i] = f(b[i-1][j-i],  
                      b[i][j-i-1], a[i][j]);  
    }  
}
```



DOACROSS Loops with OpenMP

- OpenMP 4.5 extends the notion of the ordered construct to describe loop-carried dependencies

- Syntax (C/C++):

```
#pragma omp for ordered(d) [clause[[,clause],...]  
for-loops
```

and

```
#pragma omp ordered [clause[[,clause],...]
```

where *clause* is one of the following:

```
depend(source)
```

```
depend(sink:vector)
```

- Syntax (Fortran):

```
!$omp do ordered(d) [clause[[,clause],...]  
do-loops
```

```
!$omp ordered [clause[[,clause],...]
```

- The ordered clause tells the compiler about loop-carried dependencies and their distances

```
#pragma omp parallel for ordered(2)
for (int i = 1; i < N; ++i) {
    for (int j = 1; j < M; ++j) {
        #pragma omp ordered depend(sink:i-1,j) depend(sink:i,j-1)
        b[i][j] = f(b[i-1][j],
                    b[i][j-1], a[i][j]);
    }
    #pragma omp ordered depend(source)
}
```

Example: 3D Gauss-Seidel

```
#pragma omp for ordered(2) private(j,k)
for (i = 1; i < N-1; ++i) {
    for (j = 1; j < N-1; ++j) {
        #pragma omp ordered depend(sink: i-1,j-1) depend(sink: i-1,j) \
            depend(sink: i-1,j+1) depend(sink: i,j-1)
        for (k = 1; k < N-1; ++k) {
            double tmp1 = (p[i-1][j-1][k-1] + p[i-1][j-1][k] + p[i-1][j-1][k+1]
                + p[i-1][j][k-1] + p[i-1][j][k] + p[i-1][j][k+1]
                + p[i-1][j+1][k-1] + p[i-1][j+1][k] + p[i-1][j+1][k+1]);
            double tmp2 = (p[i][j-1][k-1] + p[i][j-1][k] + p[i][j-1][k+1]
                + p[i][j][k-1] + p[i][j][k] + p[i][j][k+1]
                + p[i][j+1][k-1] + p[i][j+1][k] + p[i][j+1][k+1]);
            double tmp3 = (p[i+1][j-1][k-1] + p[i+1][j-1][k] + p[i+1][j-1][k+1]
                + p[i+1][j][k-1] + p[i+1][j][k] + p[i+1][j][k+1]
                + p[i+1][j+1][k-1] + p[i+1][j+1][k] + p[i+1][j+1][k+1]);
            p[i][j][k] = (tmp1 + tmp2 + tmp3) / 27.0;
        }
    }
    #pragma omp ordered depend(source)
}
```

OpenMP API Version 5.1

State of the Union

Architecture Review Board

The mission of the OpenMP ARB (Architecture Review Board) is to standardize directive-based multi-language high-level parallelism that is performant, productive and portable.



Development Process of the Specification

- Modifications of the OpenMP specification follow a (strict) process:



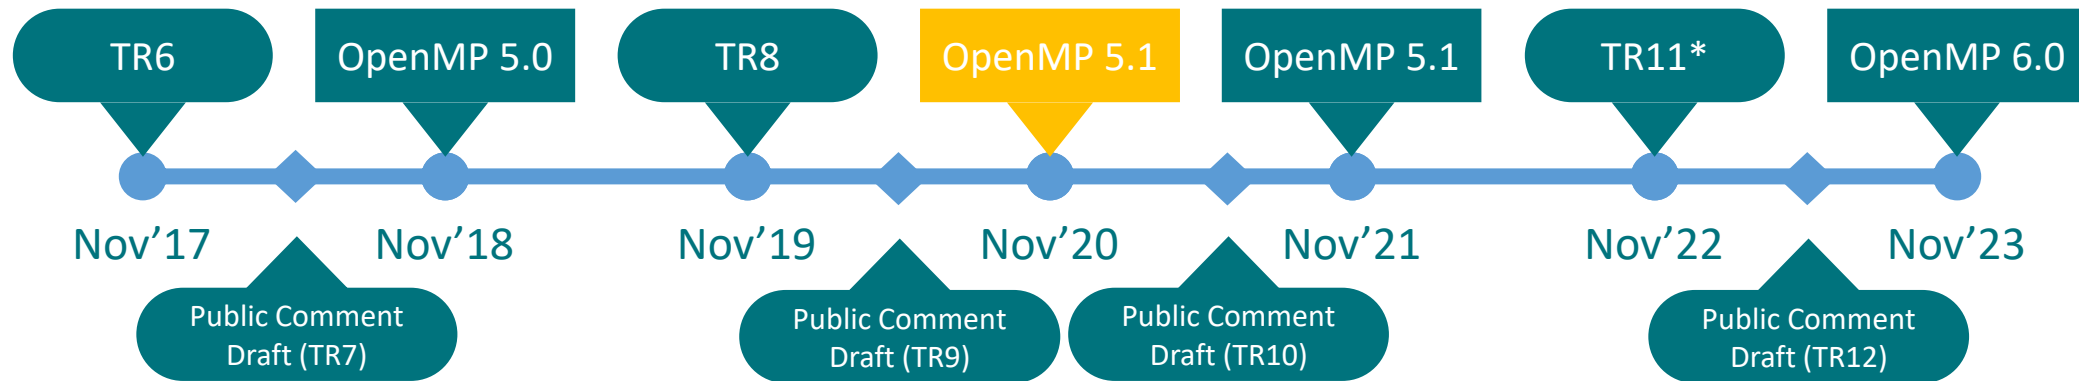
- Release process for specifications:



OpenMP Roadmap

■ OpenMP has a well-defined roadmap:

- 5-year cadence for major releases
- One minor release in between
- (At least) one Technical Report (TR) with feature previews in every year



* Numbers assigned to TRs may change if additional TRs are released.

OpenMP API Version 6.0 Outlook – Plans


- Better support for descriptive and prescriptive control
- More support for memory affinity and complex memory hierarchies
- Support for pipelining, other computation/data associations
- Continued improvements to device support
 - Extensions of deep copy support (serialize/deserialize functions)
- Task-only, unshackled or free-agent threads
- Event-driven parallelism

Printed OpenMP API Specification




- Save your printer-ink and get the full specification as a paperback book!
 - Always have the spec in easy reach.
 - Includes the entire specification with the same pagination and line numbers as the PDF.
 - Available at a near-wholesale price.
- Get yours at Amazon at <https://link.openmp.org/book51>

Recent Books about OpenMP



Covers all of the
OpenMP 4.5 features, 2017



Introduces the
OpenMP Common Core, 2019

Help Us Shape the Future of OpenMP

- OpenMP continues to grow
 - 33 members currently
- You can contribute to our annual releases
- Attend IWOMP, become a cOMPunity member
- OpenMP membership types now include less expensive memberships
 - Please get in touch with me if you are interested



Visit www.openmp.org for more information